Method	Exec.	Non-exec.	Data to Pass		Data to Reject		Metrics		
			Passed↑	False alarm↓	Rejected↑	Missed↓	Precision↑	Recall↑	F1↑
deequ	69.1	0.0	16.0	10.0	21.0	17.0	48.5% ± 0.0%	61.5% ± 0.0%	54.2% ± 0.0%
zero-shot [gemini-2.5-flash]	55.6	0.0	18.0	8.0	12.0	26.0	40.8% ± 1.2%	69.2% ± 10.9%	51.3% ± 3.9%
zero-shot [gpt-4.1]	40.6	0.0	20.5	5.5	15.0	23.0	46.8% ± 6.1%	78.8% ± 19.0%	58.6% ± 10.1%
zero-shot [gpt-5-mini]	34.2	0.0	24.5	1.5	7.5	30.5	44.5% ± 1.3%	94.2% ± 2.7%	60.5% ± 1.7%
zero-shot [gemini-2.5-pro]	40.0	0.0	18.0	8.0	17.0	21.0	46.2% ± 5.3%	69.2% ± 5.4%	55.4% ± 5.6%
zero-shot [gpt-5]	61.7	0.0	9.5	16.5	25.0	13.0	42.0% ± 5.5%	36.5% ± 8.2%	39.0% ± 7.0%
few-shot [gemini-2.5-flash]	42.5	0.0	20.0	6.0	8.0	30.0	40.0% ± 1.7%	76.9% ± 5.4%	52.6% ± 2.7%
few-shot [gpt-4.1]	36.4	0.0	23.0	3.0	5.5	32.5	41.4% ± 1.0%	88.5% ± 5.4%	56.4% ± 2.0%
few-shot [gpt-5-mini]	24.7	0.0	24.0	2.0	4.0	34.0	41.4% ± 1.0%	92.3% ± 0.0%	57.2% ± 1.0%
few-shot [gemini-2.5-pro]	19.6	0.0	25.0	1.0	8.0	30.0	45.6% ± 4.9%	96.2% ± 5.4%	61.8% ± 5.7%
few-shot [gpt-5]	26.0	0.0	23.5	2.5	12.5	25.5	48.0% ± 1.3%	90.4% ± 2.7%	62.7% ± 0.5%
swe-agent [gemini-2.5-flash]	13.6	0.0	26.0	0.0	2.0	36.0	41.9% ± 0.0%	100.0% ± 0.0%	59.1% ± 0.0%
swe-agent [gpt-5]	18.5	0.0	25.0	1.0	4.5	33.5	42.7% ± 0.2%	96.2% ± 5.4%	59.2% ± 0.9%
prismaDV [gemini-2.5-flash]	61.5	0.0	15.5	10.5	22.5	15.5	50.2% ± 2.3%	59.6% ± 8.2%	54.3% ± 2.1%
prismaDV [gpt-4.1]	59.0	0.0	15.0	11.0	28.5	9.5	61.0% ± 6.3%	57.7% ± 10.9%	59.2% ± 8.7%
prismaDV [gpt-5-mini]	63.4	0.0	19.5	6.5	28.5	9.5	67.3% ± 0.8%	75.0% ± 2.7%	70.9% ± 0.7%
prismaDV [gemini-2.5-pro]	40.9	0.0	22.0	4.0	24.0	14.0	61.2% ± 2.4%	84.6% ± 0.0%	<u>71.0% ± 1.6%</u>
prismaDV [gpt-5]	88.8	0.0	24.0	2.0	30.5	7.5	76.2% ± 1.7%	92.3% ± 0.0%	**83.5% ± 1.0%**

Table 1: Detection performance with respect to the impact of data errors on downstream tasks in EIDBench-real (overall, all 3 datasets pooled, mean ± std across runs). The best F1 (mean) is in bold, the second-best underlined. `Exec.` reports the average number of executable constraints, `Non-exec.` the average non-executable. Convention: *safe = positive class*; Passed = safe bundles correctly cleared, False alarm = safe bundles incorrectly flagged, Rejected = harmful corruptions correctly detected, Missed = harmful corruptions missed.